

LAB 1: IMAGE RECOGNITION


Due: 11:59pm on Thursday, September 7th

Individual assignment

The objective of the first lab is to use image processing and machine learning tools to correctly classify images. Specifically, you will develop an algorithm that will enable your Cozmo to distinguish the following seven symbols from each other, and from other random scenes the robot might see. Additionally, you will use RANSAC to fit a line to the arena wall.



Your algorithm will be given images taken by the Cozmo as input. Your algorithm must then determine which of the above symbols if any is in the image. You will also implement line fitting using RANSAC in order to detect arena edge lines. Note: you will not be given a physical Cozmo, but the provided images are captured from real Cozmos.



Images for classification: We have provided an image dataset containing grayscale images taken by the robot. The dataset contains 8 types of images: one for each of the above symbols, plus the type “none” for pictures not containing a symbol (just a picture of the empty arena, cube or wall). The dataset is further split into separate testing and training datasets.

Images for line fitting: We have provided an image dataset containing grayscale images of the arena wall for testing with RANSAC.

Installation: You will be developing your algorithm in [Python 3](#), using [scikit-image](#) and [scikit-learn](#). If you are already experienced in installing Python packages, feel free to ignore the following instructions:

0. Windows users: installing [Git Bash](#) first may make installation easier.
1. [Install Python 3](#) for your OS.
2. Use pip to install scikit-image and scikit-learn:

On MacOS and Linux, run this in a new terminal:

```
python3 -m pip install scikit-learn scikit-image
```

On Windows, using the [Git Bash](#) terminal:

```
py -m pip install scikit-learn scikit-image
```

Reading: Read through the “Getting Started” and “A crash course on Numpy for images” section of the [scikit-image user guide](#). Also read “[an introduction to machine learning with scikit-learn](#)” of [scikit-learn user guide](#). You may use any function from either of those libraries for your task.

Lab: We have provided the following files in canvas, under Files/Labs/Lab1_release:


train/ - directory containing images to be used for training your classifier. The correct label of each image is listed in its file name.

test/ - directory containing images to be used for testing your classifier. The correct label of each image is listed in its file name.

wall/ - directory containing images to be used for line fitting using RANSAC.

imgclassification.py - this is the main file where you will enter your solution.

The code already contains functions for reading in the images from the train/, test/ and wall/ directories and for formatting the data into an array. Add code to convert pixel values of the images to features that can be used for training a classifier. Then train a classifier of your choice and test its predictive performance against the test set. You will need to fill in three functions in the ImageClassifier class, namely, `extract_image_features`, `train_classifier`, and `predict_labels`. You can run this file locally to see your score on the training dataset.



For `line_fitting` you will need to fill in code to use RANSAC to estimate the best line (ex. red line on the left image) that fits on the wall of the arena. For reliable execution, this will run on wall images, which are images with the arena wall without the symbol cards. As mentioned in the lecture, you should run RANSAC on edges extracted from an edge detector in order to reduce noise. You're free to use any edge detector you desire, such as Canny edge detector.

Do not modify the header and the return type for each function as that will interfere with the autograder and affect your grade. You may add helper functions as needed but make sure your code is self-contained because this is the only file you will be submitting.

Evaluation: We will evaluate the performance of your algorithm on 186 images from a withheld test set containing images that were taken under the same conditions as the ones you have been provided, but that were not included in your dataset. (This is standard practice in machine learning to make sure your model avoids [overfitting](#).) The classification training set will remain the same. Your grade will be determined as the percentage of the images classified correctly.

Wall images within the withheld test set will be used for line fitting evaluation. Note that in this assignment, all tested wall images will contain a visible line indicating the boundary between the arena wall and empty space.

Submission: The assignment is due by 11:59 pm on Thursday, September 7th. You must not rename the file from `imgclassification.py` and submit it on Gradescope. Do not zip the file. Do not upload any additional files. Make sure your code is contained entirely within this file. If you relied significantly on any external resources to complete the lab, please reference these in the file comments. This assignment is to be completed individually.

Grading Rubric:

Image Recognition: 70 pts total based on percentage of images classified correctly:

$$\frac{\text{correctly classified images}}{\text{total images}} \times 70$$

15 pts out of the 70 pts for image recognition will be based on the images given to you, meaning 55 out of 70 pts are based on the hidden images.

Line Fitting: 30 pts based on percentage of lines identified correctly:

$$\frac{\text{correctly identified lines}}{\text{total images}} \times 30$$

5 pts out of the 30 pts for line fitting will be based on the images given to you, meaning 25 out of 30 pts are based on the hidden images.